

The University of Texas at Austin
Department of Integrative Biology
2415 Speedway
Austin, TX 78712

June 21, 2024

Editors
Proceedings of the National Academy of Sciences

Subject: manuscript revision

Dear Editors,

Thank you for considering this work for publication in *PNAS*. Our manuscript, “Diverse Genotype-by-Weather Interactions in Switchgrass” has undergone substantial revision from our original submission in 2021. This timeframe is long - Dr. MacQueen took a new position in 2022. They continued working with the group as we developed a new algorithm to address major reviewer concerns and significantly altered the manuscript. Because of the major alterations over this time frame, the response to reviewers does not have track changes for each point, though we do point to line numbers for the associated revisions in our response.

We were pleased by the useful feedback given by the reviewers of our original submission and believe that the greedy mash algorithm we implemented addresses the major statistical concern of reviewer 2, as well as substantially improving our use and interpretation of covariance matrices in *mash*. We have addressed all the editorial issues suggested by the two reviewers, and included in this letter are point-by-point responses to these issues.

Reviewer #1:

Suitable Quality?: No
Sufficient General Interest?: No
Conclusions Justified?: Yes
Clearly Written?: No
Procedures Described?: Yes

Comments on Significance Statement:

Conclusions are very specific for the crop/population and traits. General conclusions are lacking.

Comments:

The paper by MacQueen et al. presents results of a study on two switchgrass populations grown in 8 garden experiments across a range of latitudes. Traits are green-up date and flowering time. The data show strong GxE, effects of QTL changing in magnitude and sign. Results are partly validated in an independent segregating population.

The topic is timely and highly relevant. A better understanding of GxE will be valuable for plant genetics, evolution and breeding. The data are complex and the authors have succeeded in making them accessible to the reader.

Relevant literature

The introduction and discussion are focused on findings in Arabidopsis and switchgrass. An excellent study from Arabidopsis with high relevance for this paper is Fournier-Level et al. 2016 <https://doi.org/10.1073/pnas.1517456113>

Major findings from other crop species have been ignored. To mention just two references:

Bustos-Korts et al. 2019 The Plant Journal <https://doi.org/10.1111/tpj.14414>

Millet et al. 2016 Plant Physiology <https://doi.org/10.1104/pp.16.00621> and

There are many papers on crop modelling (eg work by Mark Cooper and others in maize) and the authors might want to consider linking their results to crop modelling to demonstrate the relevance of the results of this study

I'm still reading through these to determine which to add. May need to make additional edits to the Intro/Discussion to add these citations.

Results

There is a large overlap of data and findings with reference 32 Lovell et al. 2021 Nature. It has not become clear which additional insights can be gained from this study. One reason might be that objectives of the study have remained rather vague as well as the generic conclusions that can be drawn from the results.

While we do rely on the same genotypic data and common gardens used in Lovell *et al.*, 2021, and make use of the population structure findings of Lovell *et al.*, 2021 to divide our genotypes into subpopulations for GWAS, the previous work considered fitness proxies and this work considers phenological traits. This work has two additional, key advances over the previous work: it can assign loci genome-wide to multiple kinds of GxE and GxWeather, and has an unbiased statistical method to identify loci with and without rank-changing GxE. We now summarize the findings from Lovell *et al.*, 2021 that we use as a springboard for this work on lines 136-145, and outline the additional advances of this study on lines 156-163, 169-171, and 883-893.

In the introduction the aim is given as „we test if these populations differ in their phenological adaptation and hence their phenological GxE“. Given prior knowledge such as the origin of the populations and earlier results (Lovell *et al.*) what is expected?

We now edit the Introduction to more clearly articulate the aims of this manuscript. Our key expectation is that different genetic subpopulations and genomic regions have likely evolved distinct patterns of GxE (lines 55-58; lines 219-223). Thus, we aim to identify the kinds of GxE present for each subpopulation and each trait. (lines 129-131).

The results for these two populations did not show a general pattern leading to conclusions on the environmental cues. Associations trait/cue changed across populations, traits, gardens, analysis etc. , it was difficult to see the big picture.

We agree that we do not identify one general pattern or one environmental cue affecting any set of populations or phenological trait. Instead, our key expectation is that different populations and different genetic loci will have evolved distinct patterns of GxE (lines 55-58; lines 219-223; SI Appendix Section S1). Thus, we aim to identify the kinds of GxE present genome-wide and elaborate when these kinds of GxE differ (Figure 2), and demonstrate that individual loci have different types of GxE across environmental regions and subpopulations (Figure 3).

The authors claim that the environmental cues in the hypothesis based models improved model fit. What inference is possible from this result? Unless results are validated in independent data or with cross validation the predictive power of the cues cannot be judged.

We edit the Results to (i) change our analysis add an algorithm to more rigorously demonstrate this claim given the data that we have, and (ii) more clearly articulate this claim. We now employ a greedy mash algorithm to iteratively add covariance matrices to the *mash* model, only adding matrices that significantly improve the mash model log-likelihood, as we elaborate on further below in our [response to reviewer 2](#).

We agree that given the data we have, the predictive power of these cues cannot be judged - in fact, it is hard to imagine how our GxWeather covariances could be used predictively, as we define the GxWeather covariance matrices retroactively from each genotype's green-up date and flowering date (lines 231 - 235; SI Appendix, Section S1).

Results from the SNPs “mash model of Midwest green-up fell on a covariance matrix of average temperature in the 10 days prior to Midwest green-up”. These results should be linked to the phenotypic GxE analysis.

We edit the results to directly compare the GxWeather matrices and the phenotypic correlations in Figure 1B. (lines 241-258; lines 269-276; lines 283-287).

Flowering posterior weights were higher than green-up weights. What does that mean? How does that link to the different heritabilities and levels of GxE?

We believe this comment refers to there being higher loadings on the GxWeather matrices for flowering date than there were for the date of the start of vegetative growth in the original paper (original Figure 2). Our revised analyses uses a selection process for canonical and GxWeather covariance matrices to select only the subset that significantly improves the model log-likelihood, using a [greedy mash algorithm to iteratively select matrices](#). Thus, the covariance matrices included in each model have changed in this revision, as have the relative loadings on GxWeather and canonical covariance matrices.

Now, posterior weights on the GxWeather matrices vary between 0 to 0.6, but we do not see clear patterns between posterior weights between phenotypes or subpopulations, save that for most models, a large fraction of effects, 0.2 to 0.6, load onto a matrix of ‘No Effects’ that shrinks effects towards zero. We don’t believe it’s appropriate to make direct comparisons between heritability or a variance components analysis and these mash loadings, so we do not do so in the paper; instead of the variance due to additive genetic variation and GxE, these show the additive genetic effects that most closely matched different patterns of GxE and GxWeather, regardless of the proportion of additive variance each locus has.

What can you conclude from the pattern of antagonistic pleiotropy between Texas and Northern gardens that has not been seen in the GxE analysis?

No conclusions about individual loci effects can be drawn from the GxE analysis in Figure 2, as this plot does not show how shrinkage changes the jointly re-estimated SNP effect sizes. The GxE analysis in Figure 2 shows the types of GxE and GxWeather present genome-wide, from loadings on individual loci, but does not show the effects of these loci on traits. Rather, it is an overall characterization of GxE across all eight gardens.

The GxE analysis in Figure 3 shows the re-estimated effects of these loci on traits as contrasts between individual pairs of gardens. From Figure 3, you can conclude that many loci have effects with antagonistic pleiotropy at the site level.

Discussion

“we must understand the current patterns of trait covariation across environments,”. What is the understanding from your data? Are there any general patterns or do we need to assess them individually for different genetic material, environments, traits? It looks like the latter so what are the consequences?

We think that these patterns need to be assessed individually for different genetic material, environments, and traits (lines 220-223); thus, we develop an approach to specify multiple environmental cues and compete them to explain patterns of genetic effects (lines 230-237): we use a greedy algorithm to use with mash to select covariance patterns, and use mash to flexibly identify these patterns across populations and environments. (lines 238-240).

Need to add more here to fully address this point

“this is the first experimental work using QTL mapping and GWAS across”. I disagree, see references above and others.

We edit this to clarify our claim. NB: We do cite the first experimental work using both QTL mapping and GWAS ([Brachi et al., 2010](#)), and there are many others; we meant using these approaches to map multiple types of GxE.

“Gulf and Midwest subpopulations have two distinct photoperiod-related flowering responses....”. This has been shown for other species, eg Unterseer et al. 2016 Genome Biology

Try and add this reference in the introduction

Data/Methods

Is one year data enough to make conclusions about the effect of environmental cues? It is well known that GxYear is more pronounced than GxLocation.

We agree that this would be an interesting topic for future work. Our aim was to demonstrate the value of this approach for mapping GxE, and we believe one year of data sufficient to do this.

The reader should be able to understand what the hypothesis in the hypothesis-driven models is without looking at the supplement.

Agreed. We add Table 1 to better explain the covariance matrices we select from using the greedy mash algorithm. We include the matrices selected by this algorithm as visualizations in Figure 2.

Reviewer #2:

Suitable Quality?: No
Sufficient General Interest?: Yes
Conclusions Justified?: No
Clearly Written?: No
Procedures Described?: No

Comments:

The authors present a massive dataset on phenological variation among switchgrass cultivars across 8 common gardens. They use this dataset to describe the genetic architecture

of gene-environment interactions in two phenological traits: the timing of green-up and flowering. They find that the genetic architecture of these traits varies considerably among two populations of switchgrass (called Gulf and Midwest) and among the latitudinal cline in common gardens, and that some of this variation seems to be related to variation among gardens in key environmental cues including temperature and daylength. They also identify genomic regions associated with this variation and replicate these regions in a separate set of F2 populations.

While previous publications from this experiment have been published, I believe this is the first to focus on these two phenological traits which are of clear importance in switchgrass. The overall questions that they target are also of great importance - understanding the genetic basis of gene-environment interactions, and identifying the environmental drivers of these phenological traits - and the analytical approach that they use is creative and leverages powerful statistical methods. Given this, this study has the potential to be an important case study in the field for identifying GxE loci. However there are a number of issues with the statistical approach and also with the conceptual framework that I think make the current results uninterpretable. Also, because the methods are relatively novel and complex, it would be very helpful to provide more intuition behind the analytical approaches and more access to the raw results so that others can understand the approach more fully.

First a note: It would be helpful if the authors would provide a pdf with line-numbers and with a font that can be copied directly.

Edited. We apologize for this omission.

Major issues:

Conceptual issues:

- I like the idea of using environmental measures near the time of phenological events to try to identify environmental drivers and compare the plasticity functions across environments. However the choices here do not make physiological sense, especially for flowering. The environmental indices focus on either the day of flowering, or the 1-2 weeks prior to flowering. However, the developmental transition to flowering likely occurs well before this interval (I'm not sure how long in switchgrass, but it's likely much earlier). Photoperiod and temperature cues driving this developmental transition are irrelevant once the developmental commitment has been made. The length of time before flowers emerge and open may be dependent on these environmental factors as well, but this is physiological, not related mechanistically to the daylength or temperature requirements for flowering. The intervals used may be more relevant to green-up, I'm not familiar with the developmental basis of this trait. But as is, I don't think these metrics are interpretable for what they are designed for.

We agree that the choice of environmental measure is an extremely important consideration, and we recognize that the time window the weather signal is integrated over could be long. We wanted our model to inform us on reasonable time windows, rather than having us assert these (probably incorrectly) in our model. Initially, we thought computational feasibility constrained the number of covariance matrices we could add to the mash model, as the runtime of mash increases with the number of covariance matrices included. To address this comment and the mash statistical issue the reviewer points out below, we now specify many more GxWeather covariance matrices with more time frames - 1-7 days prior to the phenological event, and 14, 21, and 28 days prior to the phenological event, generating 48-60 matrices per weather variable. Then, we used a greedy mash algorithm, [explained below](#), to select GxWeather covariance matrices from this set that significantly improve the model likelihood. We thus extended the time frame for GxWeather covariance

that our models could capture; however, no 21 or 28 day matrices were selected by the greedy algorithm.

We also now try and explicitly describe the two phenotypes more clearly in the text - both of the traits we map were measured on individuals, but specifically when approximately half of the ramets of the genotype had open flowers (for flowering) or green leaves (for green-up). Thus we are looking for cues driving vegetative and reproductive transitions for a majority of ramets, not for heading (flower emergence on the panicle). While there may be physiological mechanisms driving flower opening once the transition to reproductive development has been made, our results show that there is also genetic variation underlying differences in timing of the physiological mechanisms underlying flower opening.

- I've struggled to understand the "hypothesis-based covariance matrix" idea. Since this idea is really key to the idea of this paper, I think the authors should work to make it much more clear and intuitive, perhaps a diagram could help. I think the idea is: if in two gardens the daylength-at-flowering metrics are correlated across accessions, this implies that the importance of daylength on flowering is similar in these two gardens. In a third environment that has little variation in daylength-at-flowering, the genetics of the daylength pathway are still the same, so the lack of observed variation must mean that this cue is less important here. If a SNP has a similar effect in these two gardens, but a smaller effect in gardens without much variation in this metric, this implies that this SNP may contribute to the signaling pathway that links daylength to flowering. This is a neat idea. However the way this is described in the text (eg "SNP effects on flowering ... covaried with daylength") is confusing. This sounds like SNP effects were larger in gardens where the mean daylength at flowering was later. Or maybe that these SNPs were also GWAS hits for daylength at flowering itself.

We thank the reviewer for this insight; we have reframed the way we talk about these covariance matrices in the paper (starting with the title, abstract, and significance statement; continuing throughout). Instead of hypothesis-based covariance matrices, we emphasize that the matrices we create are looking at an interaction between genetics and weather - they are GxWeather matrices, and we look for effects with different GxWeather interactions that are based on specific weather cues.

Statistical issues:

- Narrow-sense heritability: A key result is that the h^2 is low when measured across all trials, but high within trials. However the model used to estimate the global h^2 is missing a term for replication of lines across accessions. u only accounts for replication of additive effects of lines across gardens, but non-additive genetics can be persistent too. Also, this model doesn't allow $\text{var}(e)$ to vary across locations but it's likely that it does. If the goal is to show rank-crossing GxE, actually fit a GxE model like was done for the environmental indices to directly show it.
- Variance components analysis: The specification of this model is incorrect because $\text{Var}(u)$ and $\text{Var}(ul)$ are assigned the same distribution. There are many ways of specifying GxE effects in a model like this, and it's not clear how it is done here. If G is a $n \times n$ GRM for n accessions, it can't be the covariance matrix for both $\text{Var}(u)$ (dimension $n \times n$) and $\text{Var}(ul)$ (dimension $n_l \times n_l$). Generally, $\text{Var}(ul)$ would be $G \Psi$ where Ψ is a $l \times l$ covariance matrix among gardens. Sometimes this is constrained to be diagonal (ie no covariance among gardens), and sometimes to be proportional to I (constant variance and no covariance among gardens), but these more restrictive models should be justified.

We agree that correctly specifying the type of GxE in our models of additive variation across gardens is important for all models in this manuscript. In many cases the complexity of the models that we can specify using this data is limited, because of the low sample size - we typically have one replicate per genotype in each common garden, and the low number of genotypes present at each garden in each strongly stratified subpopulation means that we cannot fit complicated GxE structures using linear mixed models, often needing

to fit very simple/restrictive variance-covariance structures. The entire focus on *mash* was to allow us to fit less structured variance-covariance structures, as well as more of them.

Thus we thought hard about the minimum number of linear mixed models we could specify to justify our modeling approach with *mash*. We think that the presence of negative phenotypic covariation and additive genetic variation at each garden is sufficient to motivate a search for the sign- or rank-changing GxE that could potentially underlie strong negative covariation between the North and Texas gardens. Clearly, additive genetic variation is necessary to conduct a genetic mapping analysis; in addition, we find many more loci with rank-changing GxE for green-up than for flowering, and green-up has negative phenotypic correlations which flowering does not.

We could not find a similar compelling reason to include the variance components analysis of the phenological traits or weather-derived traits based on these phenological dates. Because we construct the GxWeather matrices from covariation in these weather-derived cues, then *mash* loads loci onto these GxWeather matrices, we felt it did not add value to the paper to include a variance components analysis of these as additional traits; thus, we remove the variance components analysis from this revision.

Additionally, we move the narrow-sense heritability analysis to the SI Appendix (Figure S1). Our goal was to show suitability for further analysis of additive genetic variation, not directly show rank-crossing GxE with this mixed linear model, so we do not change the environmental effect of site, which does not have an interaction term, in the narrow-sense heritability models.

- GWAS: The GWAS model the authors use is not specified (nor was it specified in the 2021 Nature paper referenced). I believe looking at the source code that the model is a linear model with PCs to account for structure but no kinship matrix. It's not clear how many PCs were used and since the raw GWAS results are not presented it's impossible to tell if this was sufficient to account for the significant structure in these populations. I'm concerned about this because of the extremely unlikely number of significant markers (19K LD blocks). If each of these was a true positive, the average effect size of each block would be only 0.005% of the genetic variance. With only 350 accessions, the power to detect a locus that explains 10% of the genetic variance is only 50% with $\alpha = 10^{-5}$ and MAF = 0.05. So if the model is detecting tons of loci it's likely that the model is severely biased by population structure, and these biases are likely to be somewhat consistent across gardens because the structure is similar. *mash* doesn't help when the individual models are biased, it'll find patterns of correlation whether they are biological or not and use those to shrink effect sizes together.

First, we now better document our GWAS methodology for this manuscript in Section S2 of the SI Appendix. We also include additional information on GWAS, including the number of PCs used and the Manhattan and QQ plots and associated data tables in the Github repository associated with the paper.

Second, though we selected 19K LD blocks in our set of 'strong' effects for *mash*, very few of these effects were significant (<500 relatively unlinked LD blocks per set of eight gardens; this still represents some genomic inflation, no doubt, but not nearly as bad as 19K significant blocks would be). We now clarify this important detail in the main manuscript and in the SI Appendix.

Third, population structure could certainly bias the resultant *mash* models. We describe the steps we take to prevent "garbage in" to our *mash* models, including removing conditions where GWAS had significant population structure. However, it is our intuition that, should biases remain in our models, as they undoubtedly do, *mash* would not load these biased effects onto our GxWeather covariance matrices - instead, we these effects could perhaps load onto canonical matrices, such as garden-specific effects, or perhaps onto data-driven matrices. Thus, we expect our assessment of the proportion of the genome that shows GxWeather effects to be conservative if there is residual population structure.

- Mash: A lot of the results interpretation is based on interpreting the SNP loadings on the specified covariance matrices. This is a secondary use of mash (the primary being the refinement of effect sizes), and while they do interpret these somewhat in the mash paper, these loadings need to be interpreted with caution. If these covariance matrices are similar, mash will somewhat arbitrarily assign weight to any of the matrices because all lead to the same fit to the effect-size data. The mash algorithm doesn't give the full posterior distribution on the loadings so you can't check for posterior correlations there. It seems likely to me that the hypothesis-driven and data-driven covariance matrices are somewhat correlated here, and the correlations may differ between green-up and flowering because of the better correlation between the environmental metrics and flowering than green-up. In the mash paper, they used cross-validation with the likelihood in the testing set as the evaluation metric to compare models. It might be safer to try dropping specific covariance matrices and comparing the model performance in a held-out testing set to evaluate the importance of the hypothesis-derived covariance matrices.

We thank the reviewer for this comment and agree that the selection of covariance matrices to be included in a mash model is a nontrivial question. To clarify our interpretation of the posterior matrix weights provided by mash, we performed an extensive analysis of the performance of mash models when different covariance matrices are included. Specifically, we implemented a model selection approach that uses a greedy algorithm to evaluate the log likelihood of the mash model as additional covariance matrices were included (see Methods lines ___ and pseudo code below). This is a similar, but more computationally efficient design, than the leave-one-out approach recommended by the reviewer as the runtime of mash increases with the number of covariance matrices included. Additionally, while cross-validation is a powerful approach to model evaluation our limited sample size was prohibitive to the necessary partitioning of individuals included in our analysis.

In practice this greedy algorithm approach offers a fast way to identify the point where redundancy in the addition of a new covariance matrix results in no change to the likelihood of the mash model. Importantly, this removes the potential for arbitrary assignment of weights to the covariance matrices in the model and allows for the selection of matrices that most accurately capture the underlying biology captured by the site-specific effect sizes. We applied this algorithm to only the hypothesis matrices for each phenotype (Figures 1 and 2) using the previously used sets of significantly associated and randomly selected variants in the Gulf subpopulation individuals, Midwest subpopulation individuals, and the two subpopulations combined. The results indicate that only a small number of the original hypothesis matrices (between three and five) are necessary to reach the maximum likelihood model when using the same likelihood ratio test implemented in Urbut et al., indicating that correlation among hypothesis covariance matrices is high. This redundancy was confirmed through our application of the greedy algorithm to a combined set of canonical, data driven, and hypothesis covariance matrices (Figures 3 and 4) Our primary conclusion from this analysis is in line with what the reviewer posited, that there is extensive redundancy among canonical, data driven, and hypothesis covariance matrices that is leading to biologically informative patterns of covariance between site effect sizes to be arbitrarily distributed among them. We additionally found evidence that while the correlations between data driven and environmentally informed (hypothesis) covariance matrices, there is extensive similarity in the matrices that are included in the maximum likelihood models across both phenotypes and in the analysis of the Gulf subpopulation, Midwest subpopulations, and the combined cohort.

Greedy algorithm for mash model selection.

```
-----  
matrices = []  
most_likely_matrices = []  
maximum_likelihoods = []  
  
likelihood = -inf  
most_likely_matrix = '
```

①
②


```

for matrix in matrices:
    matrix_likelihood = mash(effects, std.errs, matrix) ['likelihood']
    if matrix_likelihood > likelihood:
        most_likely_matrix = matrix
        likelihood = matrix_likelihood

matrices.remove(most_likely_matrix)
most_likely_matrices.append(most_likely_matrix)
maximum_likelihoods.append(likelihood)

lrt_pvalue = 0
iteration = 2
while lrt_pvalue < 0.05:
    likelihood = -inf
    most_likely_matrix = ''
    for matrix in matrices:
        model_matrices = most_likely_matrices + matrix
        model_likelihood = mash(effects, std.errs, model_matrices) ['likelihood']
        if matrix_likelihood > likelihood:
            most_likely_matrix = matrix
            likelihood = matrix_likelihood

    matrices.remove(most_likely_matrix)
    most_likely_matrices.append(most_likely_matrix)
    maximum_likelihoods.append(likelihood)

    lrt_pvalue = likelihood_ratio_test(likelihood, maximum_likelihoods[iteration-1], df = 1)

```

- ① Initialize covariance matrices of interest
- ② Initialize lists to store the most likely additional matrix and corresponding maximum likelihood
- ③ Obtain model likelihood for a model fit with each matrix
- ④ Remove the matrix with the greatest likelihood from the array and store it
- ⑤ Repeat the process of fitting mash models, finding the most likely, pair, trio, etc. (using the most likely matrices from the preceding iteration) until the likelihood ratio test is no longer significant

- There appears to be an issue with the construction of the hypothesis-derived covariance matrices. I believe the goal here is to estimate the genetic covariance among gardens based on the phenotypic covariance (P) and the estimated residual covariance (E), where $P = G + E$. The described approach involves starting with the the phenotypic correlation matrix and then replacing the diagonal with the coefficient of variation. If the diagonal had been replaced with the h^2 in each garden, that'd be a valid estimate of G. But as described, it's likely that you'll end up with an invalid covariance matrix, one that's not positive-semi-definite. Instead the appropriate operation would be to both pre- and post-multiply the covariance matrix by a diagonal matrix with the square-root of the CV's. This maintains a valid covariance-like structure. Or just mean-standardize the traits first and then calculate the covariance of this standardized matrix.

We now use the narrow-sense heritability as the diagonal for each GxWeather matrix, and clarify this in the Supplementary Methods (SI Appendix, Section S1, first and eighth paragraph - first and last paragraph on page two).

- Antagonistic Pleiotropy: Maybe I'm missing the idea of this analysis, but it doesn't seem to me that this strategy for partitioning loci between antagonistic pleiotropy and differential sensitivity would have equal power to detect each class. It seems that antagonistic pleiotropy should be a subset of differential sensitivity because the latter is defined as simply a change in magnitude. But even if it's restricted to "different magnitude but same sign", the practical definition here seems to be that antagonistic pleiotropy is detected when the signs are different and both $\text{lfsrs} < \text{threshold}$ (0.05?), while differential sensitivity additionally requires a threshold on the difference in effect sizes (0.4) with no justification for why this size was chosen. This additional criterion will necessarily make the rate of detection different. Even without this, with differential sensitivity one effect size must be much larger than the other in absolute value (which isn't the case with antagonistic pleiotropy), so differential sensitivity loci will be more likely to pass the lfsr loci than antagonistic pleiotropy loci. Note that previous literature here tried to separate antagonistic pleiotropy from conditional neutrality (ie where in one location the effect was zero), while here the contrast is with differential sensitivity where effects are non-zero in all locations. I'm not aware of previous discussions in the literature of trying to differentiate these two specific classes of GxE variants and not really sure theoretically what the importance would be.

We believe that the use of the lfsr to detect effects with rank-changing GxE is a key advance in our manuscript, and so we have made changes to the Introduction, Results, and Materials & Methods (lines 957-997; lines 1055-1056), as well as the SI Appendix, to explain and motivate this statistical change.

Our expectation is that in nature, effects will always differ. Statistically, this difference will not always be detectable. Thus, rather than ask "Are these two effects different?" - as we reasonably expect two effects to be, even if this difference cannot be measured - the local false sign rate answers a more meaningful question: Can we be confident in the sign of this effect?

We think that if evolutionary geneticists are interested in antagonistic pleiotropy and rank-changing GxE, then they should use the lfsr to measure their confidence in the sign of the effect, rather than the lfr to measure their confidence that the effect is different than zero. In the lfsr case, comparisons between antagonistic pleiotropy and conditional neutrality are no longer sensible, as we are no longer doing a statistical test that can robustly detect conditional neutrality (just as the lfr does not robustly detect antagonistic pleiotropy).

It's true that the threshold for differential sensitivity is arbitrary. We have reworded the manuscript to stress that there is equal power to detect effects of different sign as there are effects of the same sign (lines 994-997; lines 1055-1056). Another class of effects we have in our analysis, that differ from classic antagonistic pleiotropy & conditional neutrality comparisons are effects which are not distinguishable by sign or magnitude. It is not accurate that differentially sensitive loci will be more likely to pass the lfsr , because the loci are first selected by significant lfsr , then divided further into effects that can & cannot be distinguished by magnitude. Thus effects that are significant, but have small magnitude in both conditions will be labeled as not being distinguishable - essentially, these effects do not have detectable GxE. We propose that, if our intent is to detect sign-changing GxE, then we should use an unbiased statistical test to look at loci with and without sign-changing GxE, the lfsr .

Best,

Alice MacQueen and Tom Juenger